

CS 816 Software Production Engineering

# Final Project - Chat Application

**Yukta Rajapur (IMT2021066)**  
**Nilay Kamat(IMT2021096)**

**December 2024**

# Contents

<b>1</b>	<b>Introduction and Setup</b>	<b>1</b>
1.1	Introduction . . . . .	1
1.2	Objective . . . . .	1
1.3	Tools Used . . . . .	2
1.4	Setup . . . . .	3
1.4.1	Git & GitHub . . . . .	3
1.4.2	Maven . . . . .	3
1.4.3	Jenkins . . . . .	4
1.4.4	Docker . . . . .	4
1.4.5	Ansible . . . . .	5
1.4.6	Ngrok . . . . .	5
1.4.7	Kubernetes . . . . .	7
<b>2</b>	<b>Project Implementation</b>	<b>8</b>
2.1	Workflow . . . . .	8
2.2	Code Development and Source Code Management . . . . .	9
2.3	Steps to Build and Run the Project . . . . .	13
2.4	Testing . . . . .	15
2.5	CI/CD using Jenkins . . . . .	17
2.6	Containerisation . . . . .	21
2.7	Configuration Management/Deployment . . . . .	22
2.8	Kubernetes . . . . .	24
2.9	Git SCM Polling and Build Automation . . . . .	28
2.10	Using Kibana for Visualizing Application Logs . . . . .	30

2.11 Loki & Prometheus - Monitoring and Observability . . . . .	32
2.12 Working of Application & Links . . . . .	34

# Chapter 1

## Introduction and Setup

### 1.1 Introduction

Our project focuses on deploying and managing a real-time chat application built with Spring Boot for backend and ReactJS for frontend. The application enables users to send and receive messages through a WebSocket-based back-end for real-time communication, paired with a user-friendly front-end interface. The project integrates key DevOps tools such as Docker for containerization, Jenkins for automated CI/CD pipelines, Kubernetes for scalable deployment, and Ansible for configuration management. The goal of the project is to emphasize the DevOps tools and workflows used to develop, deploy, and monitor applications effectively.

### 1.2 Objective

The specific objectives of this project are:

1. Develop a full-stack application with a robust back-end and interactive front-end.
2. Set up version control with **Git** and integrate with a remote repository on **GitHub**.
3. Automate the build, test, and deployment process using **Jenkins** pipelines.
4. Containerize the application using **Docker** and manage images with a remote container registry.
5. Deploy the containerized application to a **Kubernetes** cluster for orchestration and scalability.

6. Use **Ansible** for configuration management and deployment automation.
7. Implement **testing** of backend to ensure code reliability.
8. Set up centralized **logging** and monitoring using the **Kibana for visualizations** and **Grafana, Loki and Prometheus** for enhanced observability.

### 1.3 Tools Used

1. **IntelliJ IDEA Ultimate:** A *powerful IDE and Code Editor* for Java and other languages, offering advanced development, debugging and integration features for seamless coding and deployment.
2. **Git & GitHub:** Git is a *version control system and source code management* tool that allows developers to collaborate on code and track changes. GitHub is a web-based platform for *version control and collaboration*, allowing developers to host, review, and manage code repositories using Git.
3. **Apache Maven:** A *build automation* tool that helps manage dependencies and build Java-based projects.
4. **Jenkins:** a continuous integration and continuous delivery (CI/CD) tool that automates the build, test, and deployment processes.
5. **Docker:** a *containerization platform* that enables developers to package applications and dependencies into portable, lightweight containers.
6. **Ansible:** a *configuration management tool* that automates the deployment and management of infrastructure and applications.
7. **GitHub Webhooks:** a tool that *triggers automated actions* when specific events occur in a GitHub repository.
8. **JUnit:** an open-source *testing framework* for Java that enables developers to write and run unit tests.
9. **Monitoring Stack:** A set of tools (*Prometheus, Loki, Grafana, and Kibana*) for collecting, aggregating, and visualizing metrics and logs, providing insights into application performance and health.

10. **Kubernetes:** A *container orchestration* platform that automates the deployment, scaling, and management of containerized applications.

## 1.4 Setup

We need to install and setup some tools and frameworks before getting started on the workflow.

### 1.4.1 Git & GitHub

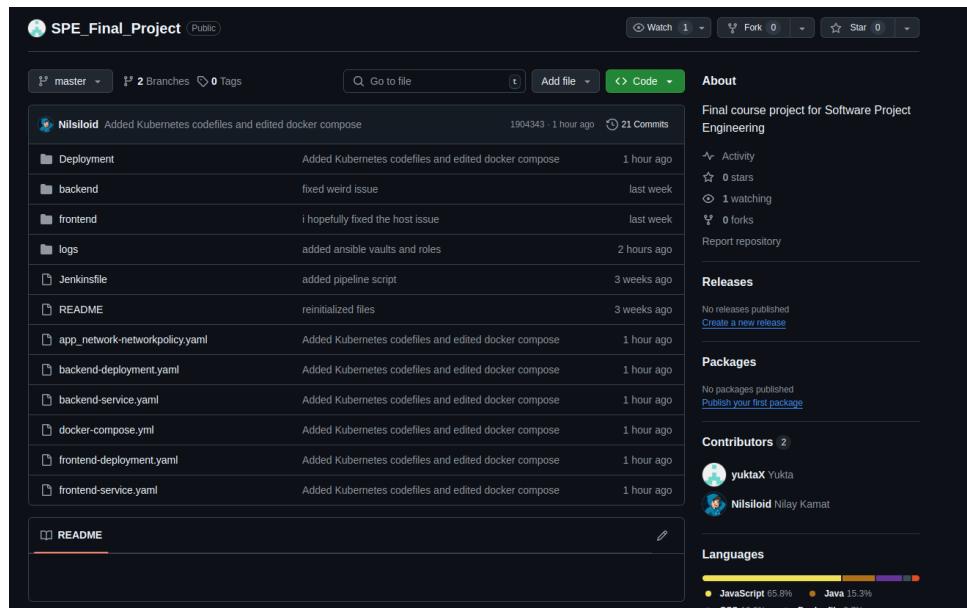
To install Git and check the version, we use the following commands:

```

1 sudo apt install git
2
3 git --version

nilay@Nilay-Lenovo-V14:~/IIITB/SEM7/SPE/Calculator$ git --version
git version 2.34.1
nilay@Nilay-Lenovo-V14:~/IIITB/SEM7/SPE/Calculator$
```

Visit GitHub and create an account if you don't have one. Then create a new repository named "SPE\_Final\_Project".



### 1.4.2 Maven

To install Maven and check the version, we use the following commands:

```
nilay@Nilay-Lenovo-V14:~/IIITB/SEM7/SPE/Calculator$ mvn --version
Apache Maven 3.9.9 (8e8579a9e76f7d015ee5ec7bfcdc97d260186937)
Maven home: /home/nilay/IIITB/SEM7/SPE/apache-maven-3.9.9
Java version: 17.0.12, vendor: Ubuntu, runtime: /usr/lib/jvm/java-17-openjdk-amd64
Default locale: en_IN, platform encoding: UTF-8
OS name: "linux", version: "6.8.0-40-generic", arch: "amd64", family: "unix"
nilay@Nilay-Lenovo-V14:~/IIITB/SEM7/SPE/Calculator$
```

### 1.4.3 Jenkins

To install Jenkins and check the version, we use the following commands:

```
● ● ●

1 sudo wget -O /usr/share/keyrings/jenkins-keyring.asc \
2   https://pkg.jenkins.io/debian-stable/jenkins.io-2023.key
3
4 echo "deb [signed-by=/usr/share/keyrings/jenkins-keyring.asc]" \
5   https://pkg.jenkins.io/debian-stable binary/ | sudo tee \
6   /etc/apt/sources.list.d/jenkins.list > /dev/null
7
8 sudo apt-get update
9
10 sudo apt-get install jenkins
```

```
● ● ●

1 sudo systemctl start jenkins
2
3 sudo systemctl status jenkins
```

### 1.4.4 Docker

To install Docker, check the version, and ensure that both Jenkins and Docker are in the same user group, we use the following commands:

```
● ● ●

1 sudo apt install curl
2 curl -fsSL https://get.docker.com -o get-docker.sh
3 sh get-docker.sh
4 sudo docker --version

nilay@Nilay-Lenovo-V14:~/IIITB/SEM7/SPE/Calculator$ sudo docker --version
[sudo] password for nilay:
Docker version 27.3.1, build ce12230
nilay@Nilay-Lenovo-V14:~/IIITB/SEM7/SPE/Calculator$ █

● ● ●

1 sudo tail /etc/gshadow
2 sudo usermod -aG docker jenkins
3 sudo systemctl restart jenkins
```

#### 1.4.5 Ansible

To install Ansible and check the version, we use the following commands:

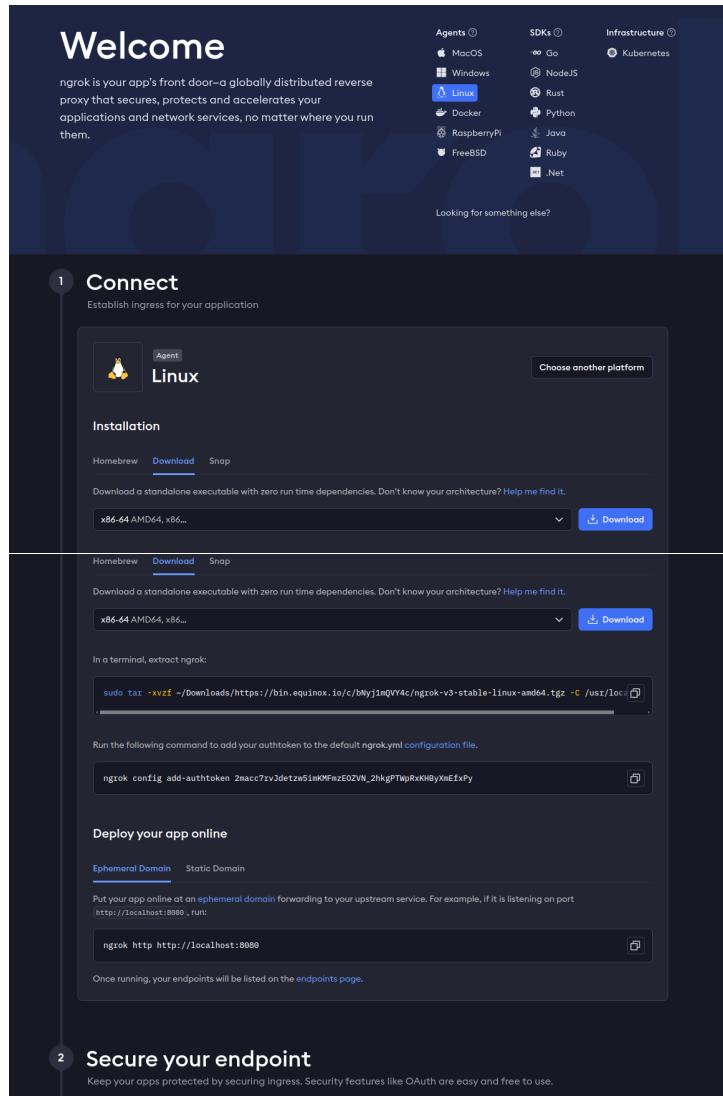
```
● ● ●

1 sudo apt-add-repository ppa:ansible/ansible
2 sudo apt update
3 sudo apt install ansible

nilay@Nilay-Lenovo-V14:~/IIITB/SEM7/SPE/Calculator$ sudo ansible --version
ansible 2.10.8
  config file = None
  configured module search path = ['root/.ansible/plugins/modules', '/usr/share/ansible/plugins/modules']
  ansible python module location = /usr/lib/python3/dist-packages/ansible
  executable location = /usr/bin/ansible
  python version = 3.10.12 (main, Sep 11 2024, 15:47:36) [GCC 11.4.0]
nilay@Nilay-Lenovo-V14:~/IIITB/SEM7/SPE/Calculator$
```

#### 1.4.6 Ngrok

Visit the Ngrok official [website](#) and sign up for a free account. Then download the Ngrok binary executable for your OS.



Once you have downloaded ngrok, you run the following command in the same directory that you have downloaded ngrok, to extract it to the destination:

```
● ● ●
1 sudo tar xvzf ~/Downloads/ngrok-v3-stable-linux-amd64.tgz -C /usr/local/bin
```

Ensure you replace ngrok-v3-stable-linux-amd64.tgz with the appropriate version of the zip you downloaded. Once you do that, you run the following command:

```
● ● ●
1 ngrok config add-authtoken 2macC7rvJdetzw5imKMFnzEOZVN_2hkgPTWpRxKHByXmEfxfPy
```

Replace the auth token with what is shown in your account dashboard.

### 1.4.7 Kubernetes

To use Kubernetes, we will install *kubectl*, the command-line tool for interacting with Kubernetes clusters, *minikube* - a lightweight tool that creates a local Kubernetes cluster on your machine.

```
● ● ●

1 sudo snap install kubectl --classic
2
3 curl -Lo minikube https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
4 chmod +x minikube
5 sudo mv minikube /usr/local/bin/
6
7 minikube start --driver=docker
8 minikube status
```

# Chapter 2

# Project Implementation

## 2.1 Workflow

1. **Code Development:** Utilised IntelliJ IDEA Ultimate and VSCode to develop the frontend and backend for the real-time chat application. The frontend was built using React.js, while the backend was developed with Spring Boot, implementing WebSocket-based communication for real-time messaging.
2. **Version Control/Source Code Management:** Stored the code in a Git repository hosted on GitHub, enabling collaborative development and version control. GitHub also facilitated continuous integration and deployment (CI/CD) workflows.
3. **Building:** Employed Apache Maven as the build tool to automate the compilation and packaging of the Spring Boot backend. For the React.js frontend, npm was used for building the production-ready application.
4. **Continuous Integration and Continuous Delivery (CI/CD):** Used Jenkins as the CI/CD tool to automatically pull code from GitHub, run unit tests, build the frontend and backend, and deploy the application to different environments. Integration with GitHub webhooks ensured that updates in the codebase triggered automated builds and deployments.
5. **Containerisation:** Docker was used to containerize both the frontend and backend applications, ensuring consistent deployment across various environments. Each service was packaged into a separate Docker image for easy scalability and deployment.

6. **Kubernetes Deployment:** Deployed the containerized applications on a Kubernetes cluster to achieve high availability and scalability. Kubernetes managed the deployment, scaling, and load balancing of the application services.
7. **Horizontal Pod Autoscaling (HPA):** Configured HPA within Kubernetes to automatically scale the backend services based on CPU usage or other metrics, ensuring optimal performance and resource allocation under varying loads.
8. **Configuration Management:** Utilised Ansible to automate the configuration of infrastructure resources. Ansible roles were used to configure Kubernetes clusters, deploy applications, and manage configurations across multiple environments, ensuring consistency in infrastructure management.
9. **Git SCM Polling and Build Automation:** Employed Ngrok to create a secure tunnel for GitHub webhooks, enabling automatic triggering of Jenkins builds whenever code is updated in the GitHub repository.

## 2.2 Code Development and Source Code Management

We have built a chat application project that allows users to communicate with each other in real-time. It is built using the Spring Boot framework for the backend, SockJS for WebSocket communication, and React for the frontend. The project utilizes Apache Maven as the build tool for the backend.

Source Code Management tools allow developers to collaborate on code, track changes and maintain version control of their codebase. They provide features such as branching and merging, which allow developers to work on different versions of the codebase and merge changes together.

### Tech Stack used:

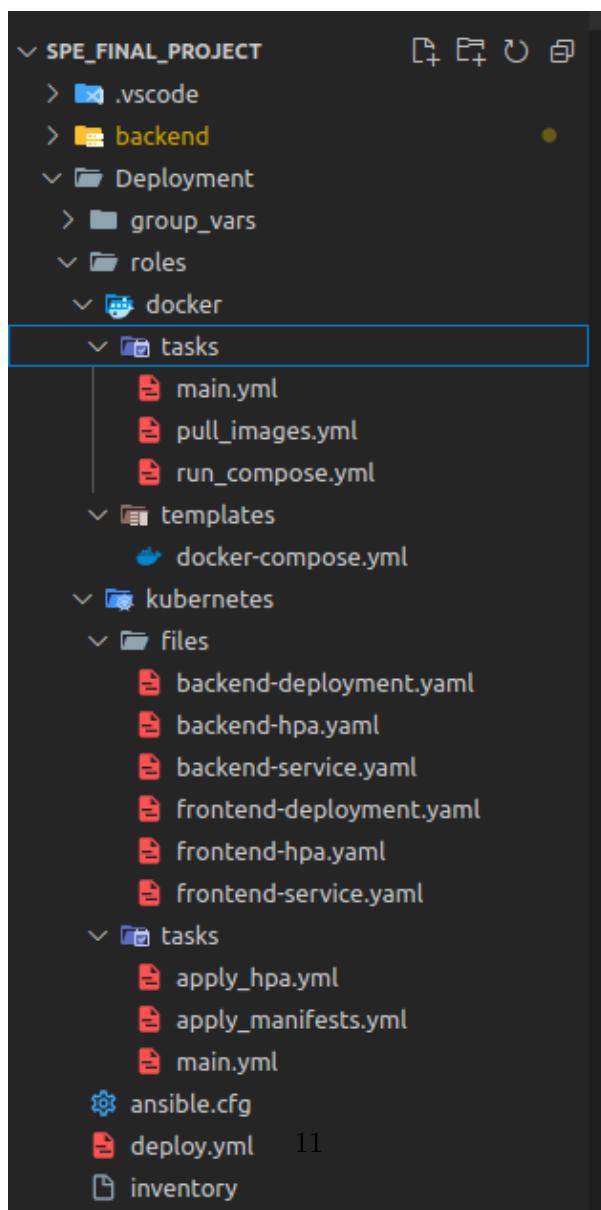
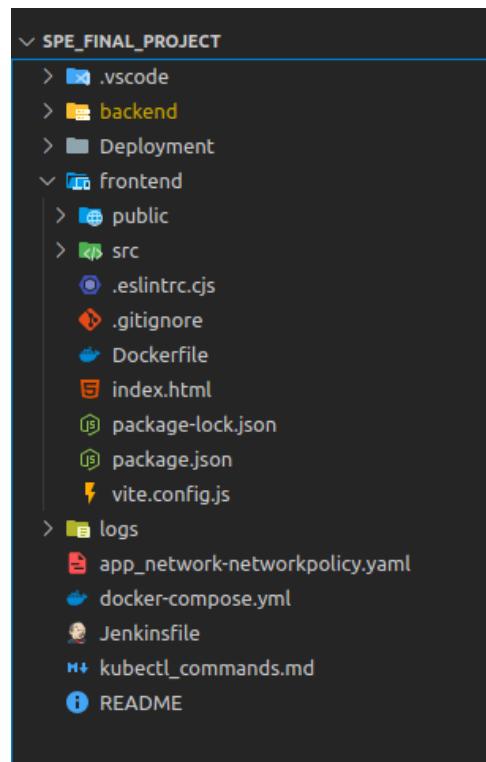
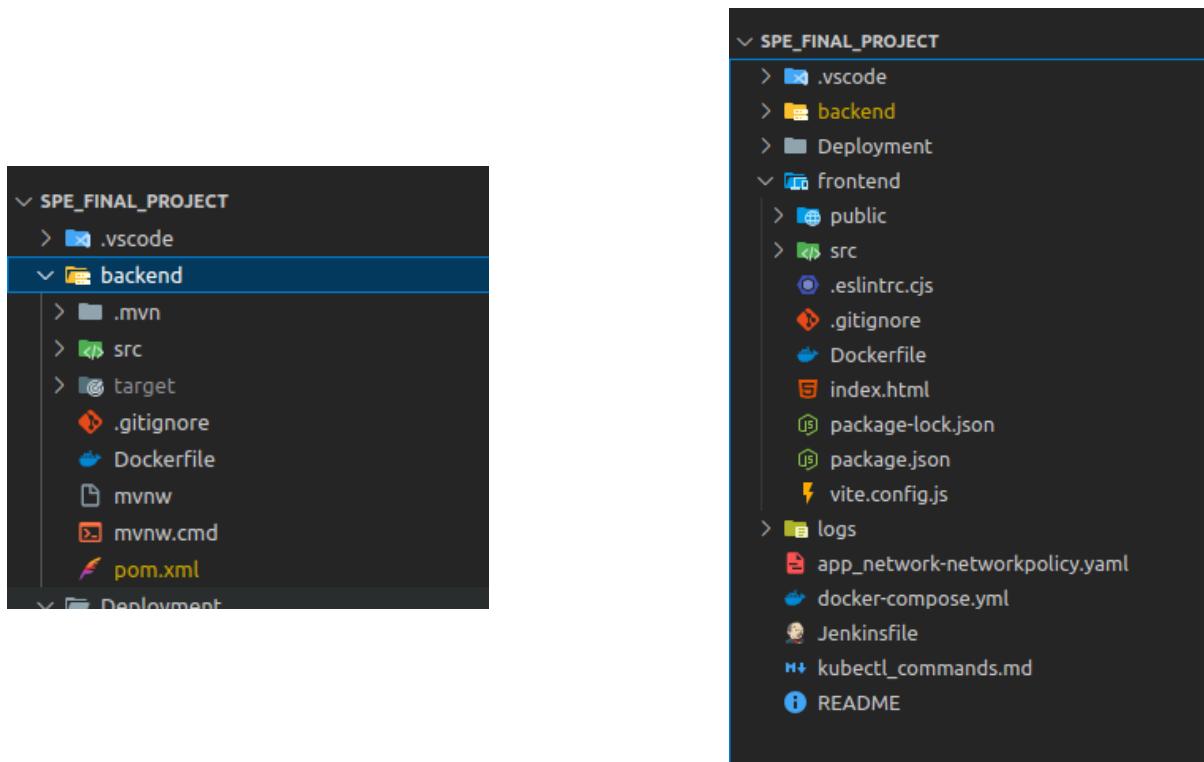
1. Spring Boot: A Java-based framework used for building the backend server and handling business logic.
2. SockJS: A WebSocket emulation library that enables real-time communication between the server and clients.
3. React: A JavaScript library used for building the user interface and handling frontend

functionality.

4. Kibana, Grafana - Tool used for Log Visualisation
5. Loki and Prometheus - Loki is a log aggregation tool, Prometheus is a metrics monitoring tool
6. Testing - JUnit, Mockito
7. Build - Maven
8. SCM - Git and GitHub

**Git Workflow and Directory Structure:** The following are the commands used to initialise a new Git repository, add files to the staging area, commit changes, set up a remote repository on GitHub, and push the changes to the remote repository.

```
1 git init
2 git add .
3 git commit -m "<message>"
4 git status
5 git remote add origin <github_repository_url>
6 git push -u origin main
7 <enter username>
8 <enter personal access token>
```

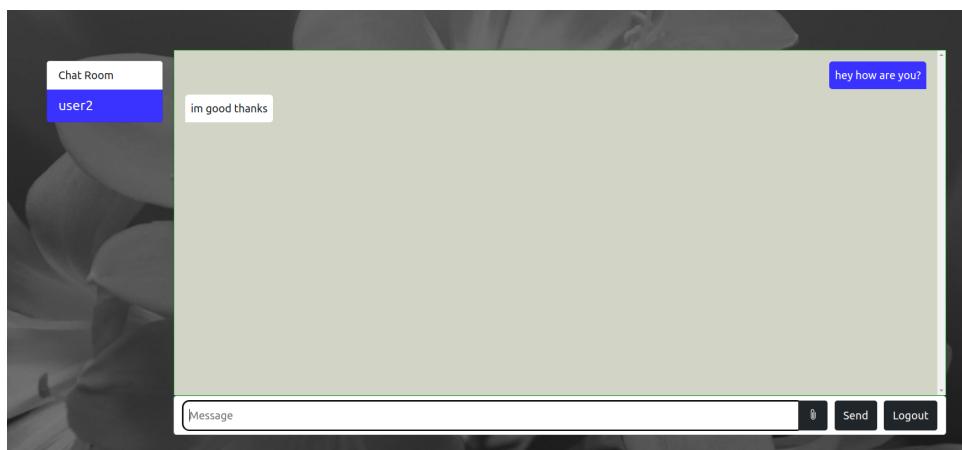


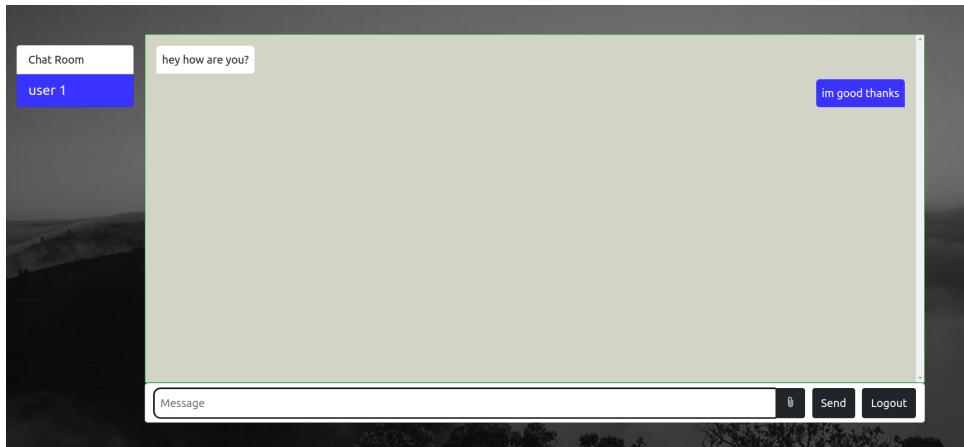
**Project Overview:****1. Folders:**

- (a) **Deployment:** Contains files for Ansible.
- (b) **backend:** Contains source code for back-end implementation done using Spring Boot.
- (c) **frontend:** Contains source code for front-end implementation done using ReactJS.

**2. Features:**

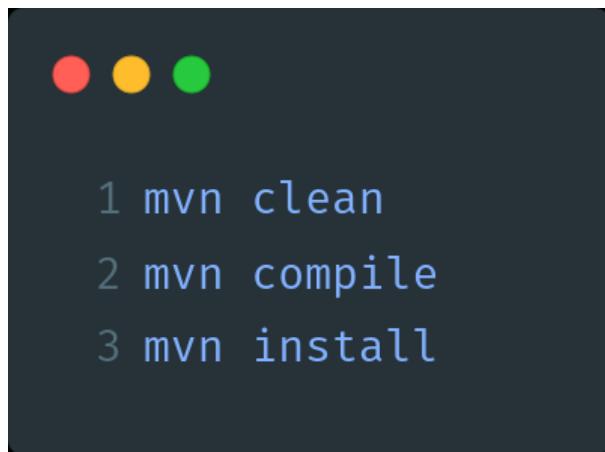
- (a) **User Authentication:** Users are required to log in with a username before they can access the chat application.
- (b) **Real-time Chat:** Once logged in, users are directed to the chat page where they can send messages in the chatroom.
- (c) **User Presence:** Users are notified when new users join the chatroom, allowing them to be aware of other participants.
- (d) **Private Messages:** Users have the ability to send private messages to specific individuals.
- (e) **Multimedia Transfer:** The application supports the transfer of multimedia files, such as photos and videos.
- (f) **Logout:** Users can log out from the application, and their username will be removed from the user list displayed to other participants





## 2.3 Steps to Build and Run the Project

To generate a JAR file with dependencies, we run the following commands:



```
1 mvn clean
2 mvn compile
3 mvn install
```

1. **mvn clean:** Removes the "target" folder, ensuring a fresh start for the subsequent compilation. This step eliminates any previous build artifacts.
2. **mvn compile:** Compiles the project and its associated test cases. This phase ensures the code is error-free and ready for the subsequent steps.
3. **mvn install:** Generates the JAR file. This final step packages the project, creating the desired output artifact (JAR file) once the compilation is successful.

Maven will build the project and check all test cases. Once completed, a "target" directory will be created in the current directory, which will contain the JAR file.

Now, we navigate to the "target" folder using the command: **cd target**.

Then, we run the JAR file using the command: **java -jar filename.jar**

Within the Maven configuration file(pom.xml), the *mainClass* tag specifies the path to the main Java file following the package structure. Additionally, the *descriptorRef* tag is employed to modify the default output JAR file name. To include project dependencies, the *dependencies* tag is utilised, enabling the addition of external libraries. In our scenario, we've incorporated dependencies such as log4j, utilised for logging functionality, and JUnit, employed for testing purposes.

#### Log4j.xml file:



```
1 %d{yyy-MM-dd HH:mm:ss.SSS}[%t] %-5level %logger{36} - %msg%n
```

The specified logging format shown above includes the following components:

1. **%dyyy-MM-dd HH:mm:ss.SSS**: Represents the timestamp with millisecond precision.
2. **%t**: Indicates the running thread's name.
3. **%-5level**: Denotes the log level with left alignment and a maximum width of 5 characters.
4. **%logger36**: Refers to the logger's name within a maximum of 36 characters.
5. **%msg**: Represents the user-written message contained in the source code.

The necessary imports for Log4j is shown in below image along with creation of an instance of the Logger in Log4j which involves using a statement to initialise a Logger object within the code.



```
1 import org.apache.logging.log4j.LogManager
2 import org.apache.logging.log4j.Logger
3
4 private static final Logger logger = LogManager.getLogger(Main.class)
```

## 2.4 Testing

The backend testing focused on validating the functionality, reliability, and performance of the chat application's core components. Using JUnit and Mockito, we ensured that each module behaves as expected in isolation and under different scenarios.

### 1. Configuration Testing (`WebSocketConfigTest.java`)

- **Objective**: To verify the WebSocket configuration for real-time communication.
- **Approach**:
  - Tested the `registerStompEndpoints` method to ensure the `/ws` endpoint is registered correctly with SockJS support.

- Verified that allowed origin patterns are set correctly to permit cross-origin requests.
  - Validated the `configureMessageBroker` method to ensure correct prefixes for application and user destinations (`/app`, `/user`).
- **Key Tools:** Mockito to mock the `StompEndpointRegistry`.

## 2. Controller Testing (`ChatControllerTest.java`)

- **Objective:** To validate the behavior of public and private messaging endpoints.
  - **Approach:**
    - Mocked the `SimpMessagingTemplate` to test private message delivery without a real broker.
    - For `receiveMessage`, ensured the controller correctly processes and returns public messages.
    - For `privateMessage`, verified that the correct user and destination are used in the `convertAndSendToUser` method.
- **Key Tools:** Mockito for mocking dependencies and `verify` for validating method calls.

## 3. Model Testing (`MessageTest.java`)

- **Objective:** To confirm the integrity of the `Message` class, which represents the data structure for messages.
  - **Approach:**
    - Conducted unit tests for getter and setter methods to ensure proper field initialization and retrieval.
    - Verified `equals`, `hashCode`, and `toString` methods for accurate comparisons and debugging outputs.
- **Key Tools:** JUnit assertions to validate object behavior.

## Testing Strategy

- **Unit Testing:** Focused on individual methods and classes in isolation to ensure correctness.
- **Mocking Dependencies:** Used Mockito to simulate real dependencies, such as the messaging template and WebSocket configuration registry.
- **Behavioral Verification:** Ensured expected interactions between components, especially in the controller.

This structured testing approach ensured that all critical backend functionalities of the chat application were thoroughly validated before deployment.

## 2.5 CI/CD using Jenkins

- Jenkins is an open-source automation tool written in Java with plugins built for continuous integration.
- Jenkins is utilized to continuously build and test software projects, simplifying the process for developers to integrate changes, and allowing users to obtain up-to-date builds with ease.
- After installing Jenkins using the procedure mentioned in the previous chapter, go to <https://localhost:8080>, then browse to **Manage Jenkins** → **Plugin Manager** → **Available Plugins** and install the following necessary plugins.
  - Git plugins & GitHub plugins
  - Maven Integration
  - Docker plugin & Docker pipeline
  - Ansible plugin
  - Kubernetes plugin
  - JUnit plugin

Following this, browse to **Manage Jenkins** → **Credentials** → **System** → **Global Credentials** and create 2 credentials as shown in image below:

The screenshot shows two separate configurations for Jenkins credentials. The first entry is for Docker Hub, with fields: Scope (Global), Username (nilay95), Password (Concealed), ID (DockerHubCred), and Description (Docker Hub Credential). The second entry is for Localhost, with fields: Scope (Global), Username (nilay), Password (Concealed), ID (localhost), and Description (Localhost User Login Credentials). Both entries have a 'Save' button at the bottom.

We will then establish a Jenkins pipeline comprising of 6 distinct stages:

1. **Stage 1 : Git Clone** - This stage clones the repository from the main branch of the provided GitHub URL.

```
1 stage('Stage 1: Git Clone') {
2     steps{
3         git branch: 'master',
4         url:'https://github.com/yuktaX/SPE_Final_Project'
5     }
6 }
```

2. **Stage 2 : Maven Build** - Executes the mvn clean install command to build the project and resolve dependencies using Maven.

```
● ● ●

1 stage('Stage 2: Setup Backend'){
2     steps{
3         sh '''
4             cd backend
5             mvn clean install
6             '''
7     }
8 }
```

3. Stage 3 : Test Backend - Executes mvn test command to test backend of application.

```
● ● ●

1 stage('Stage 3: Test Backend'){
2     steps{
3         sh '''
4             cd backend
5             mvn test
6             '''
7     }
8 }
```

4. Stage 4 : Build and Push Backend Docker image - Uses the Docker build process to create a Backend Docker image from the project. Authenticates using DockerHub credentials (DockerHubCred) and pushes the created image to Docker Hub.

```
1 stage('Stage 4: Build and Push Backend Docker Image') {
2     steps {
3         script {
4             def backendImage = docker.build(env.BACKEND_IMAGE_NAME, './backend')
5             docker.withRegistry('', 'DockerHubCred') {
6                 backendImage.push('latest')
7             }
8         }
9     }
10 }
```

5. **Stage 5 : Build and Push Frontend Docker image** - Uses the Docker build process to create a Frontend Docker image from the project. Authenticates using DockerHub credentials (DockerHubCred) and pushes the created image to Docker Hub.

```
1 stage('Stage 5: Build and Push Frontend Docker Image') {
2     steps {
3         script {
4             def frontendImage = docker.build(env.FRONTEND_IMAGE_NAME, './frontend')
5             docker.withRegistry('', 'DockerHubCred') {
6                 frontendImage.push('latest')
7             }
8         }
9     }
10 }
```

6. **Stage 6 : Clean Docker images** - Removes any stopped Docker containers and unused Docker images to free up space.

```
1 stage('Stage 6: Clean Docker Images') {
2     steps {
3         script {
4             sh 'docker container prune -f'
5             sh 'docker image prune -f'
6         }
7     }
8 }
```

7. **Stage 7 : Ansible Deployment** - Executes an Ansible playbook (Deployment/de-

ploy.yml) to automate the deployment of the application using the specified inventory file and settings.



```
1 stage('Stage 7: Ansible Deployment'){
2     steps
3     {
4         ansiblePlaybook becomeUser: null,
5         colorized: true,
6         credentialsId: 'localhost',
7         disableHostKeyChecking: true,
8         installation: 'Ansible',
9         inventory: 'Deployment/inventory',
10        playbook: 'Deployment/deploy.yml',
11        sudoUser: null
12    }
13 }
```

## 2.6 Containerisation

- Containerization is a software deployment process that bundles an application's code with all the files and libraries it needs to run on any infrastructure.
- Containers are lightweight, portable, and self-contained environments that enable developers to package an application with all its dependencies, libraries, and configuration files, ensuring that it runs consistently across different environments.
- We will use Docker to create containers. Docker is an open-source tool that enables developers to build, package, and deploy applications in a containerized environment.

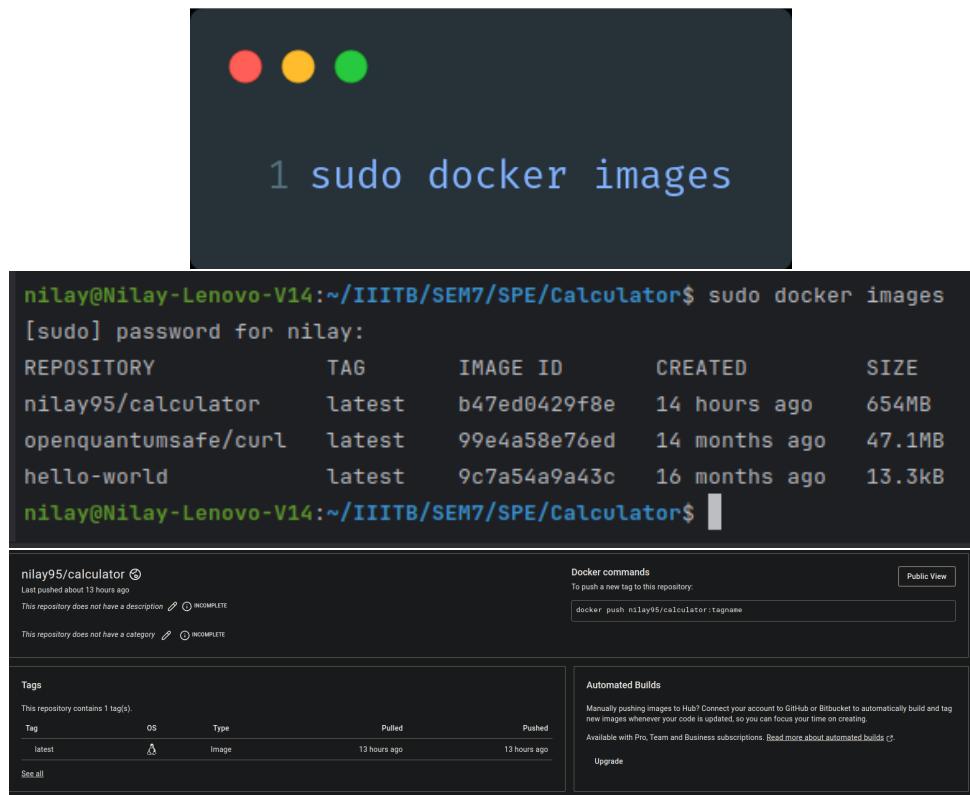
We will now create a Docker container for our project by creating a Dockerfile.

- **FROM:** We use the OpenJDK 11 base image to build our image.
- **COPY:** Copy jar file from source on the host machine into the container's file system.

- **WORKDIR:** Changes the current working directory.
- **ENTRYPOINT:** Specify the command that should be run when a container based on the image is started.

```
1 FROM openjdk:11
2 COPY ./target/Calculator-1.0-SNAPSHOT-jar-with-dependencies.jar .
3 WORKDIR .
4 CMD ["java", "-cp", "Calculator-1.0-SNAPSHOT-jar-with-dependencies.jar", "org.example.Main"]
```

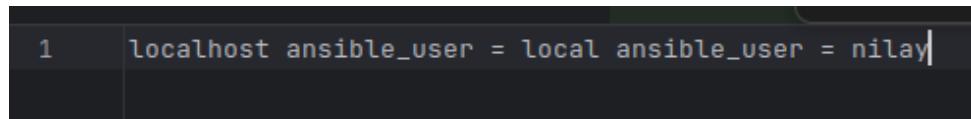
The purpose of the above code is to generate a Docker container with OpenJDK version 11 as the base image which acts like a JVM. The final container can be run independently without any dependencies. The Jenkins Pipeline script has been configured to include the Dockerfile and Docker commands, automating the process of building and pushing the Docker Image to DockerHub. In the pipeline script, docker\_image = docker.build “nilay95/calculator” builds the docker image. The following command allows verification of the successful creation of a Docker image.



## 2.7 Configuration Management/Deployment

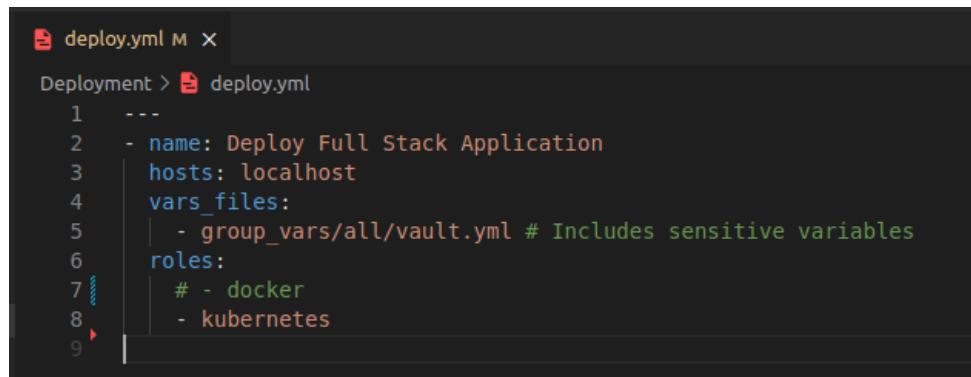
- We will use Ansible for local deployment. Ansible is a suite of software tools that enables infrastructure as code.

- In Ansible, managed hosts or servers which are controlled by the Ansible control node are defined in a host inventory file. The Ansible inventory file defines the hosts and groups of hosts upon which commands, modules, and tasks in a playbook operate.
- We are going to pull the images from the DockerHub and create containers using Ansible. We will create a Deployment folder and create two files named inventory and deploy.yml.



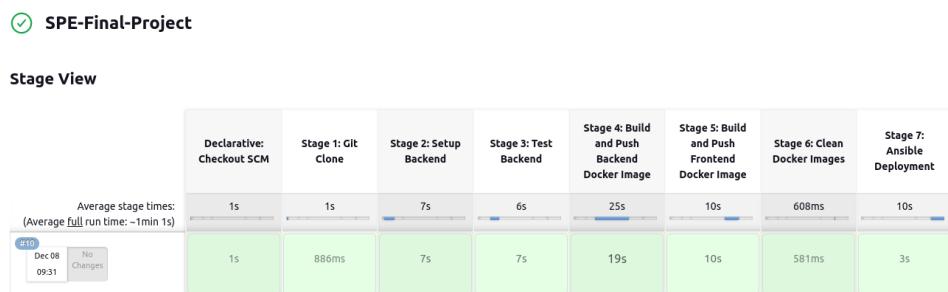
```
1 localhost ansible_user = local ansible_user = nilay|
```

- Ansible Playbooks offer a repeatable, re-usable, simple configuration management.
- Playbooks consist of one or more plays run in a particular order. A play is an ordered set of tasks run against hosts chosen from your inventory. Plays define the work to be done. Each play contains a set of hosts to configure, and a list of tasks to be executed.



```
deploy.yml M X
Deployment > deploy.yml
1 ---
2   - name: Deploy Full Stack Application
3     hosts: localhost
4     vars_files:
5       - group_vars/all/vault.yml # Includes sensitive variables
6     roles:
7       # - docker
8       - kubernetes
9 |
```

- The Jenkins Pipeline script has been configured to execute the Ansible Playbook automatically.
- At the end of these 7 steps of the workflow, the Jenkins Pipeline is completed and the stage view should be as follows:



## 2.8 Kubernetes

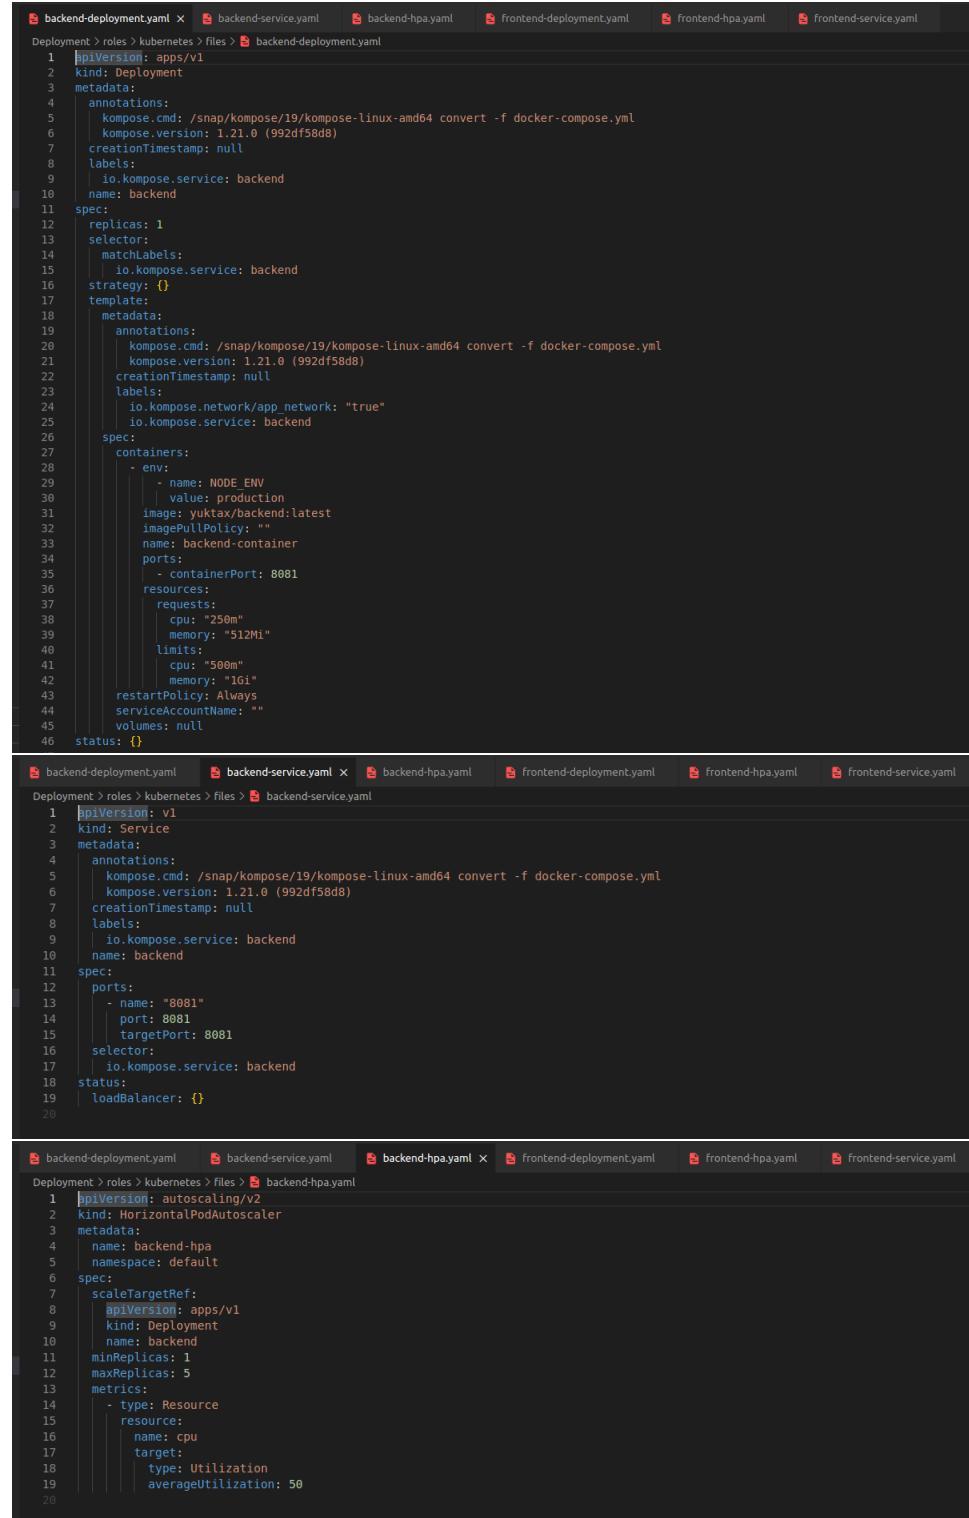
Kubernetes, commonly abbreviated as K8s, is an open-source container orchestration platform that automates the deployment, scaling, and management of containerized applications. In this project, Kubernetes was used to manage the backend and frontend services of the application, ensuring high availability, scalability, and efficient resource usage.

- **Containerization of Application:** Before deploying to Kubernetes, the application was containerized using Docker:
  1. **Backend:** A Spring Boot application was containerized as yuktax/backend:latest.
  2. **Frontend:** A React application was containerized as yuktax/frontend:latest. These images were built using Docker and pushed to Docker Hub to make them accessible to Kubernetes.
- **Setup:** Kubernetes was set up on the system to orchestrate the containerized application:
  1. **kubectl:** The Kubernetes command-line tool was installed and configured to interact with the cluster.
  2. **Cluster:** A local Kubernetes cluster was used (Minikube) to deploy the application.
- **Resources:** The following Kubernetes resources were created and applied to manage the application:
  1. **Deployments:** The backend-deployment.yaml and frontend-deployment.yaml files defined the desired state of the backend and frontend applications. Each deployment specified:
    - (a) Container image: The Docker image for the respective service.
    - (b) Replicas: The number of instances of each pod to run.
    - (c) Environment variables: Configurations for inter-service communication.
    - (d) Resource Requests and Limits: CPU and memory allocations to optimize resource usage.
  2. **Services:** Kubernetes services exposed the deployments to enable communication:
    - (a) ClusterIP for backend: Used for internal communication within the cluster.

- (b) NodePort for frontend: Exposed the frontend to the host system for external access.
3. **Horizontal Pod Autoscaler (HPA):** HPA was implemented to scale pods dynamically based on resource usage (e.g., CPU). The frontend-hpa.yaml configured the minimum and maximum number of pods and the CPU utilization target.

The files are shown in screenshots below:

- Backend Yaml files:



```

backend-deployment.yaml | backend-service.yaml | backend-hpa.yaml | frontend-deployment.yaml | frontend-hpa.yaml | Frontend-service.yaml
Deployment > roles > kubernetes > files > backend-deployment.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    annotations:
5      kompose.cmd: /snap/kompose/19/kompose-linux-amd64 convert -f docker-compose.yml
6      kompose.version: 1.21.0 (992df58d8)
7      creationTimestamp: null
8    labels:
9      io.kompose.service: backend
10   name: backend
11  spec:
12    replicas: 1
13    selector:
14      matchLabels:
15        io.kompose.service: backend
16    strategy: {}
17    template:
18      metadata:
19        annotations:
20          kompose.cmd: /snap/kompose/19/kompose-linux-amd64 convert -f docker-compose.yml
21          kompose.version: 1.21.0 (992df58d8)
22          creationTimestamp: null
23        labels:
24          io.kompose.network/app_network: "true"
25          io.kompose.service: backend
26      spec:
27        containers:
28          - env:
29            - name: NODE_ENV
30              value: production
31            image: yuktax/backend:latest
32            imagePullPolicy: ""
33            name: backend-container
34            ports:
35              - containerPort: 8081
36            resources:
37              requests:
38                cpu: "250m"
39                memory: "512Mi"
40              limits:
41                cpu: "500m"
42                memory: "1Gi"
43            restartPolicy: Always
44            serviceAccountName: ""
45            volumes: null
46        status: {}

```

```

backend-deployment.yaml | backend-service.yaml | backend-hpa.yaml | frontend-deployment.yaml | frontend-hpa.yaml | Frontend-service.yaml
Deployment > roles > kubernetes > files > backend-service.yaml
1  apiVersion: v1
2  kind: Service
3  metadata:
4  annotations:
5    kompose.cmd: /snap/kompose/19/kompose-linux-amd64 convert -f docker-compose.yml
6    kompose.version: 1.21.0 (992df58d8)
7    creationTimestamp: null
8  labels:
9    io.kompose.service: backend
10   name: backend
11  spec:
12    ports:
13      - name: "8081"
14        port: 8081
15        targetPort: 8081
16    selector:
17      io.kompose.service: backend
18    status:
19      loadBalancer: {}

```

```

backend-deployment.yaml | backend-service.yaml | backend-hpa.yaml | frontend-deployment.yaml | frontend-hpa.yaml | Frontend-service.yaml
Deployment > roles > kubernetes > files > backend-hpa.yaml
1  apiVersion: autoscaling/v2
2  kind: HorizontalPodAutoscaler
3  metadata:
4    name: backend-hpa
5    namespace: default
6  spec:
7    scaleTargetRef:
8      apiVersion: apps/v1
9      kind: Deployment
10     name: backend
11    minReplicas: 1
12    maxReplicas: 5
13    metrics:
14      - type: Resource
15        resource:
16          name: cpu
17          target:
18            type: Utilization
19            averageUtilization: 50
20

```

- Frontend Yaml files:

```

backend-deployment.yaml
Deployment > roles > kubernetes > files > backend-deployment.yaml
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4 annotations:
5 kompose.cmd: /snap/kompose/19/kompose-linux-amd64 convert -f docker-compose.yml
6 kompose.version: 1.21.0 (992df58d8)
7 creationTimestamp: null
8 labels:
9 io.kompose.service: frontend
10 name: frontend
11 spec:
12 replicas: 1
13 selector:
14 matchLabels:
15 io.kompose.service: frontend
16 strategy: {}
17 template:
18 metadata:
19 annotations:
20 kompose.cmd: /snap/kompose/19/kompose-linux-amd64 convert -f docker-compose.yml
21 kompose.version: 1.21.0 (992df58d8)
22 creationTimestamp: null
23 labels:
24 io.kompose.network/app_network: "true"
25 io.kompose.service: frontend
26 spec:
27 containers:
28 - env:
29   - name: REACT_APP_BACKEND_URL
30     value: http://backend:8081
31   image: yuktax/frontend:latest
32   imagePullPolicy: ""
33   name: frontend-container
34   ports:
35     - containerPort: 3000
36   resources:
37     requests:
38       cpu: "250m"
39       memory: "512Mi"
40     limits:
41       cpu: "500m"
42       memory: "1Gi"
43   restartPolicy: Always
44   serviceAccountName: ""
45   volumes: null
46 status: {}
47

frontend-service.yaml
Deployment > roles > kubernetes > files > frontend-service.yaml
1 apiVersion: v1
2 kind: Service
3 metadata:
4 annotations:
5 kompose.cmd: /snap/kompose/19/kompose-linux-amd64 convert -f docker-compose.yml
6 kompose.version: 1.21.0 (992df58d8)
7 creationTimestamp: null
8 labels:
9 io.kompose.service: frontend
10 name: frontend
11 spec:
12 ports:
13   - name: "3000"
14     port: 3000
15     targetPort: 3000
16 selector:
17   io.kompose.service: frontend
18 status:
19 loadBalancer: {}

frontend-hpa.yaml
Deployment > roles > kubernetes > files > frontend-hpa.yaml
1 apiVersion: autoscaling/v2
2 kind: HorizontalPodAutoscaler
3 metadata:
4   name: frontend-hpa
5   namespace: default
6 spec:
7   scaleTargetRef:
8     apiVersion: apps/v1
9     kind: Deployment
10    name: frontend
11   minReplicas: 1
12   maxReplicas: 5
13   metrics:
14     - type: Resource
15       resource:
16         name: cpu
17         target:
18           type: Utilization
19           averageUtilization: 50
20

```

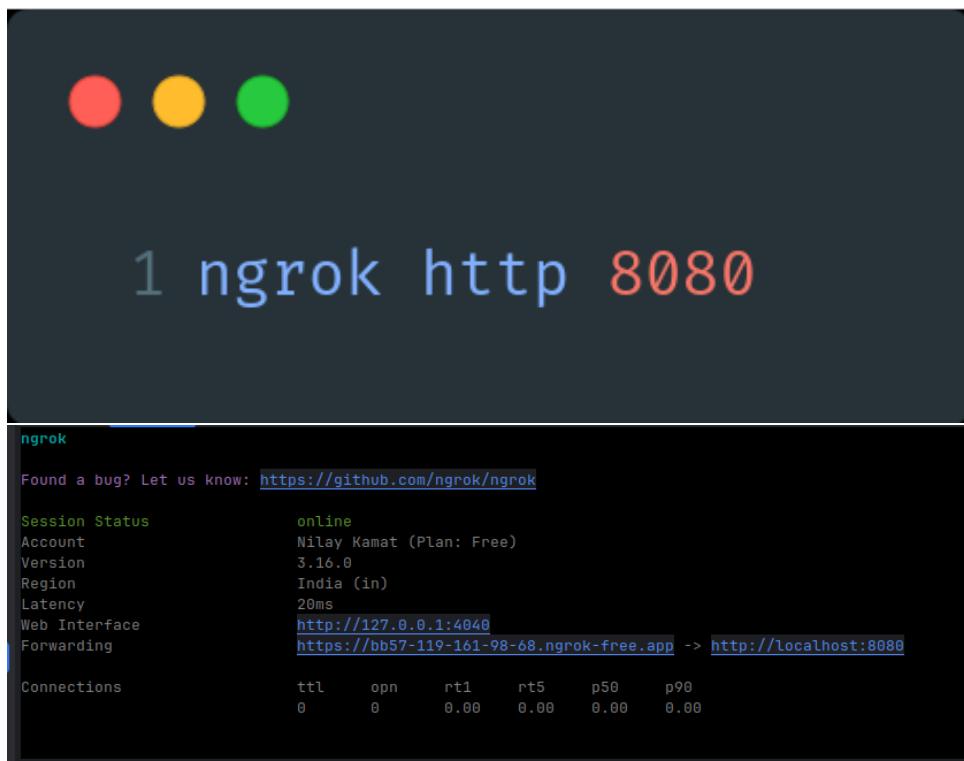
- Deployment Process:

1. **Build and Push Docker Images:** Jenkins pipelines were used to build and push the Docker images to Docker Hub.

2. **Apply Kubernetes Manifests:** The kubectl apply -f command was used to deploy the resources defined in the YAML files.
3. **Monitor and Verify:** The kubectl get pods and kubectl get services commands verified that the pods and services were running as expected.

## 2.9 Git SCM Polling and Build Automation

- Open a terminal window and enter the below command. This command will establish an HTTP tunnel using Ngrok, exposing the local server running on port 8080 to the internet.



A screenshot of a terminal window showing the output of the ngrok command. The top part of the terminal shows three colored dots (red, yellow, green) indicating the status of the tunnel. Below that, the command '1 ngrok http 8080' is displayed. The bottom part of the terminal shows the ngrok configuration details, including session status, account information, version, region, latency, web interface URL (<http://127.0.0.1:4040>), and a forwarding URL (<https://bb57-119-161-98-68.ngrok-free.app>) which maps to <http://localhost:8080>. There is also a table for connections showing ttl, opn, rt1, rt5, p50, and p90 values.

```
ngrok
Found a bug? Let us know: https://github.com/ngrok/ngrok
Session Status      online
Account            Nilay Kamat (Plan: Free)
Version             3.16.0
Region              India (in)
Latency             20ms
Web Interface       http://127.0.0.1:4040
Forwarding          https://bb57-119-161-98-68.ngrok-free.app -> http://localhost:8080
Connections
  ttl     opn      rt1      rt5      p50      p90
    0       0     0.00     0.00     0.00     0.00
```

- Copy the forwarding URL provided by Ngrok. Subsequently, create a GitHub webhook and utilise this URL as the payload URL for the webhook configuration. GitHub initiates a test connection, and upon successful configuration, a '200 OK' message confirms the proper setup.

```

ngrok
Sign up to try new private endpoints https://ngrok.com/new-features-update?ref=private

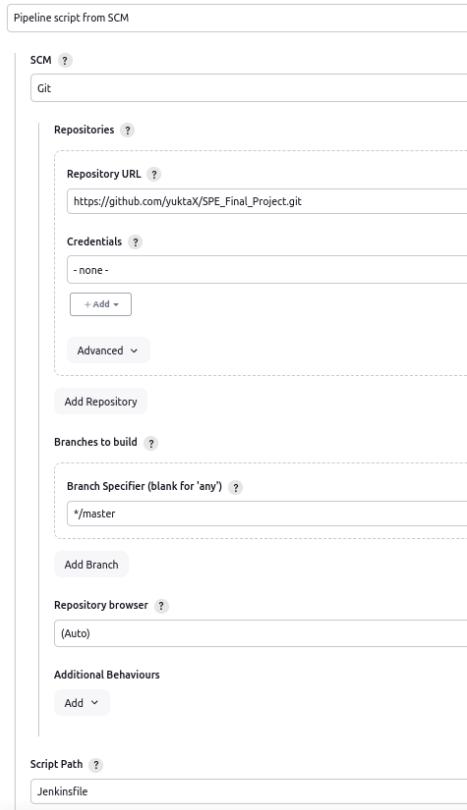
Session Status      online
Account            Nilay Kamat (Plan: Free)
Version             3.16.0
Region              India (in)
Latency             26ms
Web Interface      http://127.0.0.1:4040
Forwarding          https://6b94-103-156-19-229.ngrok-free.app -> http://localhost:8080

Connections        ttl     opn     rt1     rt5     p50     p90
                    0       1       0.00   0.00   0.00   0.00

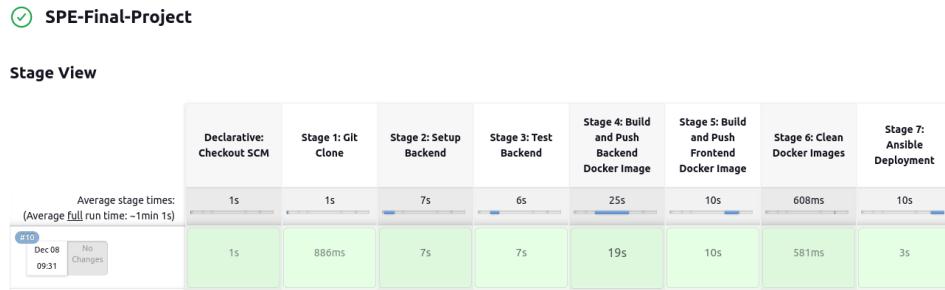
HTTP Requests
-----
11:34:03.344 IST POST /github-webhook/          200 OK

```

- Now, we'll update the Jenkins Pipeline script to a Jenkinsfile configure a build trigger for Git SCM polling. This setup ensures that our pipeline automatically initiates the build process whenever Jenkins detects a new commit made to the associated GitHub repository.



- Upon making any commits, the Jenkins Pipeline automatically initiates the build process and the final pipeline+build is as follows:



## 2.10 Using Kibana for Visualizing Application Logs

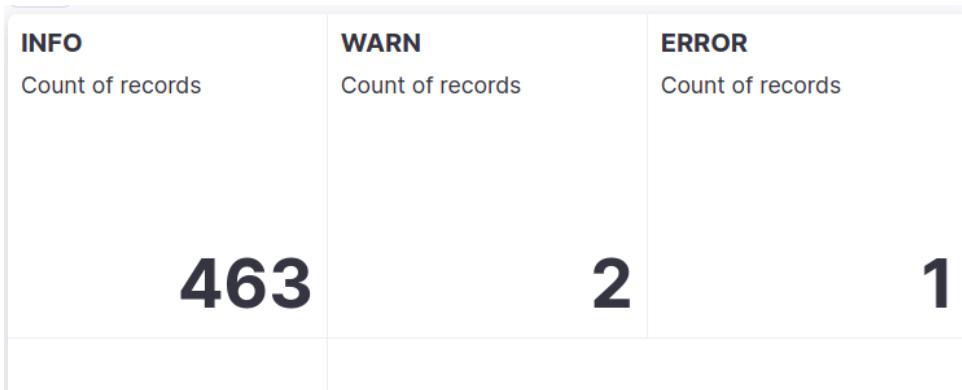
Kibana is a powerful data visualization and exploration tool that works seamlessly with Elasticsearch, enabling users to analyze and monitor application logs efficiently. By leveraging its capabilities, developers and system administrators can gain real-time insights into system performance, user activity, and potential issues.

### Steps for Visualizing Logs in Kibana

The following steps outline how to use Kibana for visualizing application logs:

- Log Ingestion:** Import log data into Elasticsearch using Logstash or Beats. For example, application logs in Grok-compatible format can be parsed and structured using Logstash pipelines. Here we simply uploaded the log file generated
- Index Creation:** Define an index pattern in Kibana that matches the imported log data. This pattern helps Kibana identify and query the logs.
- Building Visualizations:** Use Kibana's visualization tools to create charts such as line graphs, bar charts, pie charts, and heatmaps. For instance, logs with user activity can be visualized as time-series data to monitor peak usage hours.
- Dashboard Design:** Combine multiple visualizations into a cohesive dashboard for an overarching view of application behavior. For example, display logs related to error rates, active users, and WebSocket sessions on the same dashboard.



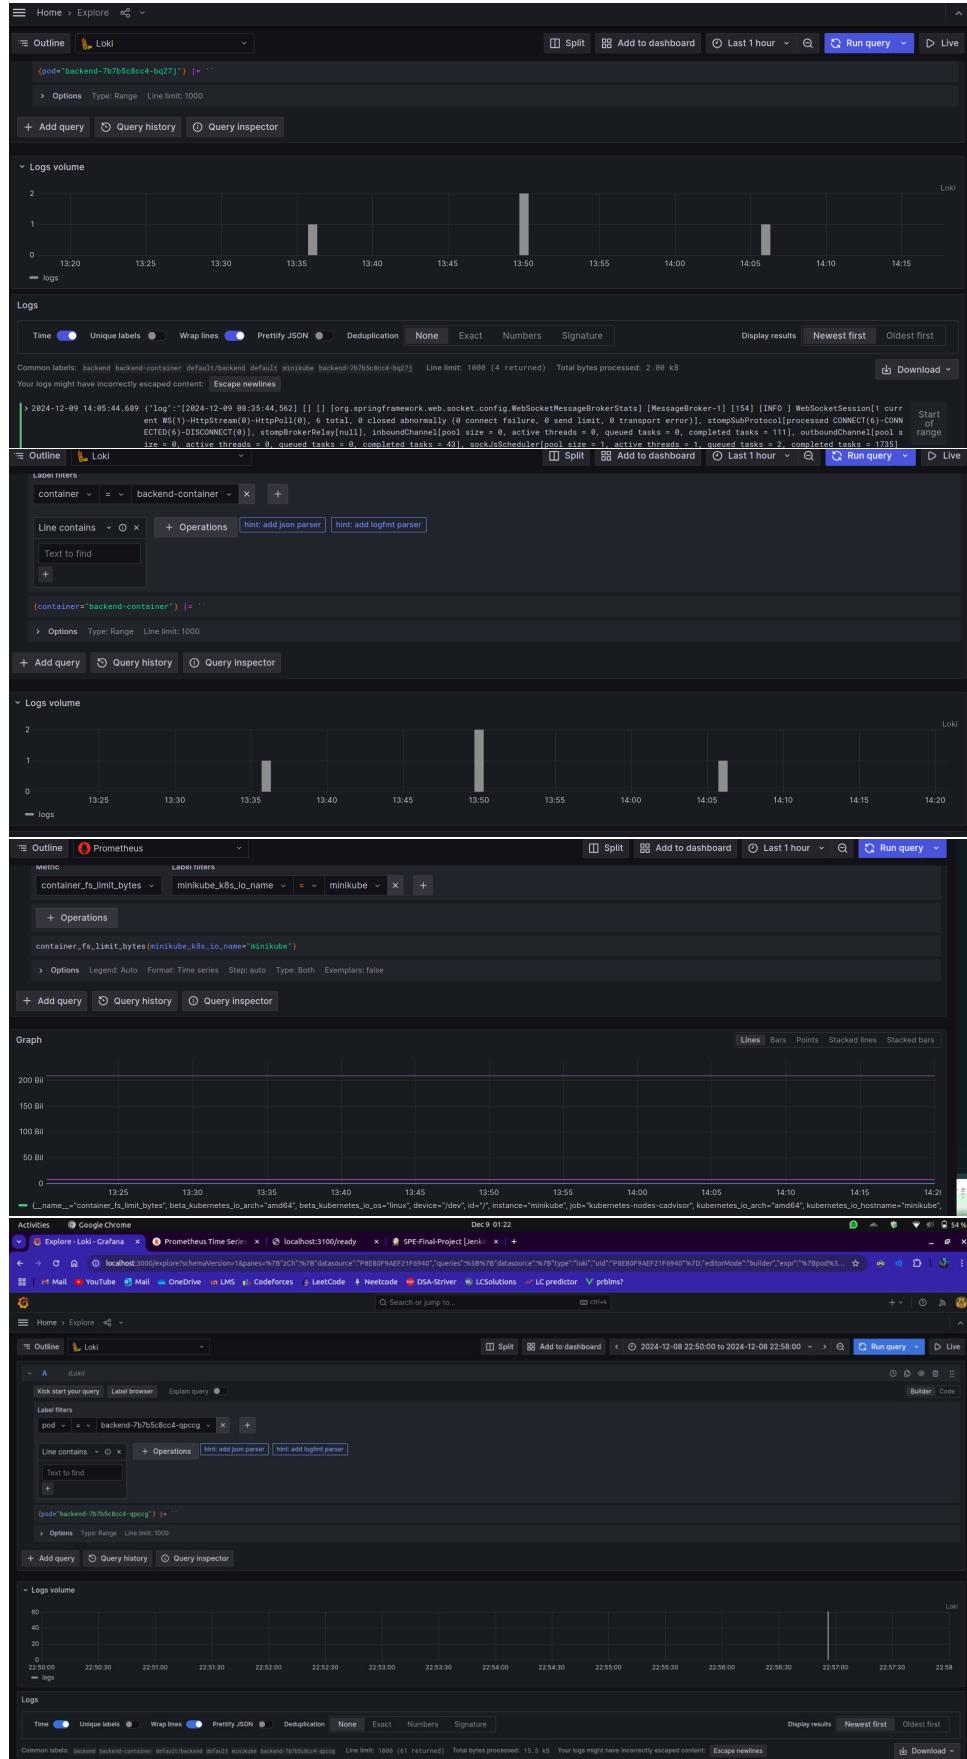


By integrating Kibana into the application's logging workflow, developers can streamline operations, ensure system reliability, and enhance user satisfaction.

## 2.11 Loki & Prometheus - Monitoring and Observability

To enable effective monitoring and observability of the application, we integrated **Prometheus**, **Loki**, and **Grafana** to interact with the k8s cluster where the application is running. The following steps outline the process. To find the exact commands used - Setup

1. **Setup Prometheus:** Prometheus was deployed to scrape metrics from the backend and frontend services as well as Kubernetes resources. Custom configurations were added to monitor CPU usage, memory consumption, and request latencies.
2. **Configure Loki:** Loki was deployed for centralized log aggregation. Promtail agents were installed to forward logs from application pods and nodes to Loki, tagging them with metadata for easy querying.
3. **Integrate Grafana:** Grafana was set up with Prometheus and Loki as data sources. Pre-built dashboards were configured for application metrics, while custom dashboards were created to visualize log patterns and performance trends.
4. **Testing and Validation:** The setup was validated by generating synthetic traffic and ensuring metrics and logs appeared correctly in Grafana dashboards. Alerts were configured for critical metrics to enable proactive monitoring.



This monitoring stack provided a comprehensive observability solution for the deployed chat

application, ensuring proactive issue detection and performance optimization.

## 2.12 Working of Application & Links

The screenshot shows a terminal window at the top displaying the command `minikube service list` and its output:

NAMESPACE	NAME	TARGET PORT	URL
default	backend	8081/8081	http://172.17.0.2:30008
default	frontend	3000/3000	http://172.17.0.2:30416
default	kubernetes	No node port	kubernetes
kube-system	kube-dns	No node port	kube-dns
monitoring	loki	No node port	loki
monitoring	loki-alertmanager	No node port	loki-alertmanager
monitoring	loki-alertmanager-headless	No node port	loki-alertmanager-headless
monitoring	loki-grafana	No node port	loki-grafana
monitoring	loki-headless	No node port	loki-headless
monitoring	loki-kube-state-metrics	No node port	loki-kube-state-metrics
monitoring	loki-memberlist	No node port	loki-memberlist
monitoring	loki-prometheus-node-exporter	No node port	loki-prometheus-node-exporter
monitoring	loki-prometheus-pushgateway	No node port	loki-prometheus-pushgateway
monitoring	loki-prometheus-server	No node port	loki-prometheus-server

Below the terminal are three browser windows, each showing a different chat room interface. All three windows have a background image of a leopard.

- Top Window:** Chat Room "test1". It shows messages "hi" and "how r u" from the user, and a reply "hi" from the server.
- Middle Window:** Chat Room "test1". It shows messages "hi" and "how r u" from the user, and a reply "hi" from the server.
- Bottom Window:** Chat Room "test2". It shows a single message "hi" from the user.

1. GitHub repository - [SPE-Final-Project Repository](#)
2. Docker images and repository - [Frontend and Backend Images](#)